

# Project Report

## IBM: Employee attrition's analysis

Edoardo Gervasoni s1043824 - edoardo.gervasoni@student.ru.nl

Alberto Monaco s1043826 - alberto.monaco@student.ru.nl

Marco Spanò s1045892 - marco.spano@student.ru.nl

Bayesian Networks course — Radboud Universiteit — January 29, 2020

### Abstract

Over the course of the first assignment we focused on the understanding of the main causes behind employees' attrition. The Bayesian network, constructed to investigate the causal relationships among the available nodes, has been built up thanks to the software *Dagitty*. After several adjustments, it has been tested through several useful methods to prove the conditional independencies returned from the network arrangement on the data. These tests showed some unexpected results (factors referring to the promotions were not very helpful to capture Attrition) and corroborated some insight previously taken into account from our *a priori* knowledge.

In this report we are going to describe the second part of this project. It will consist in applying two different algorithms, able to infer the structure of a graphical model entirely from a set of data, and in comparing their outcomes with the results previously achieved by our own Bayesian network. Throughout the project, it has to be considered the different approach while comparing the models: the Network we have built takes into account patterns and structures that come from the knowledge of the authors, their backgrounds and thoughts about the casual relationships among the nodes; the two algorithms, **PC stable** and **Tabu Search**, conversely, construct DAGs working directly on the data, without any biased judgment.

## 1 Introduction

Automatic learning for Bayesian Network is a very powerful task that is raising importance nowadays, since it permits to discover simple and easy-to-read structures from data given as input. Our work aimed to learn a Bayesian Network model from a dataset about employees' attrition and to compare it to our hand-crafted network we previously worked on.

### 1.1 Algorithms

In order to achieve this task, we decided to perform the learning of the Bayesian network using two notable algorithms belonging to two different families, PC-stable and Tabu search. Both are a better and more complex variant of two simpler algorithms from which, as we will discuss later in this section, they inherit the same approach: these are respectively *PC* for PC-stable and *hill climbing* for Tabu search.

We decided to utilize those two algorithms first because they are well-known and easy to use (they are both already implemented in several packages), second for their suitability for our dataset, and last because they approach the same problem from two really different ways, as one is a constraint-based algorithm while the other is a score-based one. In this context, it is important to compare both the results obtained by these algorithms with each other, as well as with the network that we constructed in the Assignment 1. We refer to PC-stable as the one first described in the paper of Colombo and Maathuis *A modification of the PC algorithm yielding order-independent skeletons* [3], while for Tabu search we can cite the paper *Tabu search—part I* [6] and *Tabu search—part II* [7].

PC-stable is an order-independent constraint-based algorithm that uses conditional independence tests on the dataset to discover a causal structure. This type of algorithm starts with a complete undirected graph and then prunes it by deleting all the edges that do not satisfy the conditional independence test (that could be information based, ad-hoc or statistical), in order to obtain the final skeleton of the graph. Subsequently it uses some specific rules trying to find the direction of the edges.

PC-stable is a better variant of the basic PC-algorithm, that is order-dependent and was found to be unstable especially for high-dimensional data; for example in [4], after being executed 25 times on the same dataset, it was found to have 2000 unstable edges on a total amount of 5000 in at least 50% of the cases tried. PC-stable follows the same passages of the PC but with some modifications applied, obtaining an order-independent algorithm that turns out to be more reliable and partially parallelizable.

The second algorithm is Tabu search, an iterative technique to solve an optimization problem that eventually gives us the Bayesian network as outcome. The first and simplest algorithm of this family is called *hill climbing* and it uses a score function to evaluate the actual network given a dataset (this function is the one used for the optimization problem). Starting from an initial solution (usually an empty graph), it attempts to improve the accuracy going step by step towards a local minimum.

Tabu search[2] is a memory-based strategy that changes hill climbing to end up with a better outcome. The main difference of this algorithm is that, at each step it prohibits to undo and reverse recent moves, since those would take the algorithm away from the local minimum and back to the previous state. To fulfill this requirement, the algorithm, after each step, creates Tabu entries. They consist of a list of moves that will be inserted in the table for a small number of times and that are checked at every step to block any attempt to perform them. The algorithm stops when it can't improve its current solution. Using this approach, Tabu search can outperform hill climbing, exploring more the set of possible solutions at every step (it does not go back to previous states because of the table, so it can only explore new possible paths) finding, eventually, a better one.

## 1.2 Parameters

In the two algorithms there are few parameters that can be changed, affecting the final outcome. Starting with PC-stable, two are the most important parameters that can be tuned: *alpha* and *type*.

Alpha, the parameter we chose to use in our project, is the significance level of the statistical test for conditional and unconditional independencies, and it is normally set as 0.05 as default[5]. High value of alpha will put a stricter threshold for the test of the independencies, leading the Bayesian network to obtain more edges (since the variables, according to the test, are not independent), while low values of alpha will create as output a Bayesian network with less edges.

The *type* parameter, instead, rules the variety of statistical test to perform. Different types of tests are present and usually each one is more suited to a specific kind of data.

Looking at Tabu search, two are the most important parameters: the *Score Function* and the *Tabu Length*.

The second one, Tabu length, is the parameter that governs the size of the table. As we have mentioned before, a bigger list for storing Tabu moves will push the algorithm to explore portions of the search space that were not already been visited before, improving the overall final solution. On the other hand, a bigger table will raise significantly the computational time of the algorithm, as for each step more checks of the moves must be performed. A possible default value for this parameter is 10.

On the contrary, it is possible to vary the score function of the current Bayesian Network, by changing the parameter of the same name. Several score functions are available and all of them aim to maximize the likelihood of our model looking at the data.

## 1.3 Dataset

Throughout our work we are going to use a dataset involving the attrition burden among the employees of a company[1]. The data is composed of 1470 rows and 18 columns, with all the attributes not changed from the first assignment we have worked on.

In each row of the dataset, an employee is described through those 18 variables, that represent personal (as Age, Gender, ...) or working characteristics (like Income or JobRole).

## 2 Methods

We carried out our project thanks to the R software. In particular, we have used two functions of the package *bnlearn* in order to construct the DAGs: *pc.stable* and *tabu*. The first one allowed us to build a Bayesian Network starting from the data on employee attrition's through the pc algorithm; the second one was used to learn the structure of the DAG with the tabu search approach.

### 2.1 Preprocessing

The initial part of our project consisted of a preprocessing phase, which is necessary in order to obtain data that can be processed more easily and more efficiently by the algorithms. One common problem in structure learning is the fact that usually the data available contain variables with too many levels, thus implying a high number of degrees of freedom and too many combinations of values to consider. To avoid these complications, we first selected the variables that have been used in the Assignment 1 and then we discretized those that were continuous. After that, we proceeded gradually decreasing the number of classes that the variables could take. At the end of this process, we ended up with most of our variables having 2 levels (for *Age*, *Attrition*, *BusinessTravel*, *DistanceFromHome*, *Gender*, *JobLevel*, *MaritalStatus*, *NumCompaniesWorked*, *OverTime*, *StockOptionLevel*, *TotalWorkingYears*, *WorkLifeBalance*, *YearsAtCompany*, *TotalSatisfaction*); the levels of *Education*, *MonthlyIncome* and *PercentSalaryHike* have been reduced to 3, while those of *YearsSinceLastPromotion* to 5.

### 2.2 Distance Measures

After the preprocessing step and the execution of the two structure learning algorithms, we compared the results obtained with each other and with the DAG that we manually built in the first Assignment. In order to do this, we decided to use some functions always provided by the *bnlearn* package:

- *all.equal*: checks whether two networks have the same structure.
- *hamming*: is the Hamming Distance between two networks, which computes the number of different edges in their skeletons.
- *shd*: Structural Hamming Distance, which counts how many edges are different between the CPDAGs of the two networks, taking into account their v-structures. Thus, the SHD of two DAGs belonging to the same Markov equivalence class, is guaranteed to be zero.
- *compare*: compares two networks, taking into account also arc directions, with more complete approach. The possible output are: True positive (TP), which tells us the number of arcs that are the same in both the two arguments; False positive (FP) is the number of arcs which are present in the first network (considered as baseline) but missing in the second one or that have different orientations; False negative (FN) refers to edges that are present in the second argument but not in the first one, or have different directions.
- *graphviz.compare*: is a graphical comparison from the *Rgraphviz* package, which takes one network as reference, and plots all other networks such that true positive arcs are in black, false positive arcs are in red, while false negative arcs are in blue, and drawn using a dashed line.

### 3 Results

In this section we are going to explicitly show the main differences emerged in the graphical model representations.

We will start showing the structure learning results of the algorithms Tabu Search and PC Stable, in comparison with the original Bayesian Network:

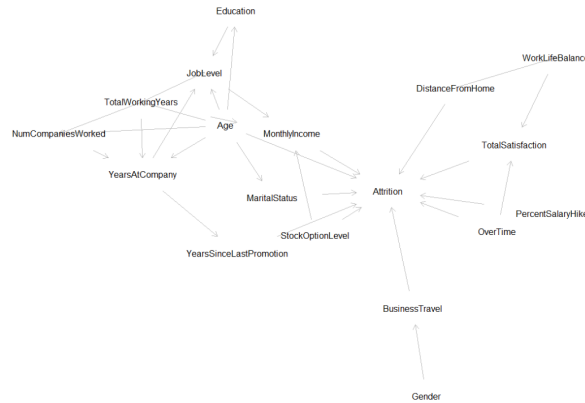


Figure 1: Network of the First Assignment

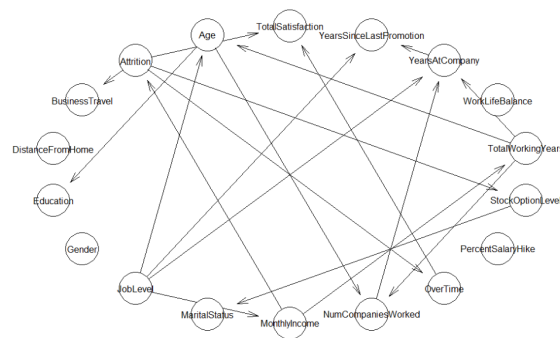


Figure 2: Network obtained with Tabu Search

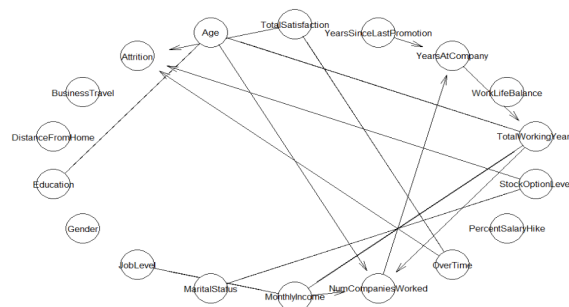


Figure 3: Network obtained with PC.stable

We investigated the differences among the networks through the indicators described in the previous section. First of all, the function *all.equal* showed that the results were all different, as we expected. Let's now look at the *Hamming Distance* and the *SHD*:

Comparison	Hamming Distance	SHD
PC - Tabu	6	13
PC - Hand-constructed	19	27
Hand-constructed - Tabu	15	32

As we can see, the major differences come from the comparison between the hand-constructed Bayesian Network with respect to the two algorithm implementations.

Now let's consider the results of the function *compare*:

First Argument v Second Argument	TP	FP	FN
PC - Tabu	3	16	12
Hand-constructed - PC	6	9	24
Tabu - Hand-constructed	9	21	10

Finally we show the differences between the DAGs obtained with *graphviz.compare*, where the leftmost Network is considered as the baseline:




Figure 4: From left to right: our hand-constructed network; its comparison with the result of Tabu Search; its comparison with the result of PC.Stable (TP in black, FP in red and FN in blue)

Once we have achieved the outcomes of the two structure learning models, we decided to vary some parameters in both the algorithms. In the first case we modified the parameter *alpha*, with levels 0.2, 0.5 and 0.75, while in the second one we varied *tabu* (the tabu table size), assigning values of 20, 50 and 100 to the table length. Here we show some of the networks obtained:

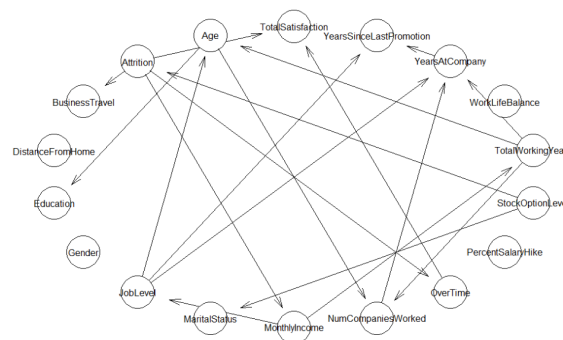


Figure 5: Tabu Search, tabu parameter: 100

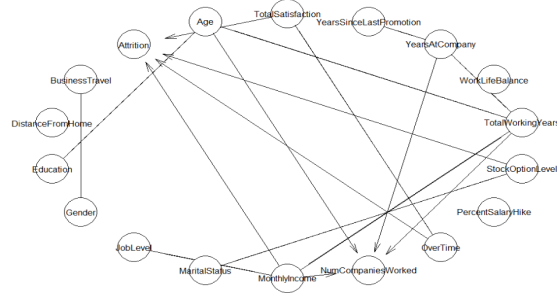


Figure 6: PC Stable,  $\alpha = 0.20$

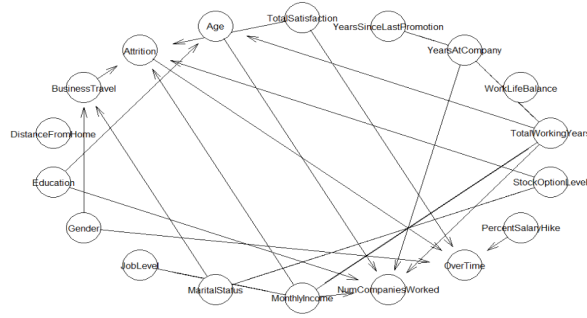


Figure 7: PC Stable,  $\alpha = 0.75$

The Tabu table size seems to not really affect the edges arrangement, except for some arrowheads (we also tried to assign lower values, obtaining the same results). On the other hand,  $\alpha$  leads the PC Stable to detect a greater number of connections in the graph. In particular, the number of edges increases at the same pace of the  $\alpha$  rise.

In the following tables, we can confirm the graphical insights received. For the PC stable:

Comparison with alpha 0.05	Hamming Distance	SHD
$\alpha = 0.2$	1	4
$\alpha = 0.5$	8	12
$\alpha = 0.75$	7	14

For the Tabu Search:

Comparison with table size 10	Hamming Distance	SHD
table size: 20	0	0
table size: 50	0	0
table size: 100	0	0

## 4 Discussion

The results showed that the two networks constructed using structure learning algorithms are less different with each other than to our manually built DAG, as we can see from the tables of the Hamming Distance and the SHD. We expected this result, given that the two algorithms consider only the data that are provided to them, while the manual construction process of the network takes into account previous

knowledge too. For instance, in our DAG we decided to keep the relations  $Education \rightarrow JobLevel$  and  $PercentSalaryHike \rightarrow Attrition$ , despite having bad results according to the data available (high p-values and low estimate coefficients in the polychoric matrix), since we thought there was a strong causal relation between these variables in the real world. The two structure learning algorithms, on the other hand, didn't maintain these edges. This discrepancy of building methods therefore inevitably leads to differences between the constructed networks.

Regarding the two algorithms, Tabu Search is the one that performed better, since it managed to detect more relationships between the variables than the PC Stable, and notably, 9 of these relationships are also present in our DAG of the first Assignment (TP = 9).

Furthermore, we noticed that, in the case of this first technique, the variation of the parameter *tabu* does not lead to significant different results, since only the direction of some edge changes; therefore we considered the default value 10 to be optimal. The statement about PC is totally different, as increasing the value of *alpha* leads to a greater number of connections detected. However, this result is due to the fact that we are just increasing the significance level of the statistical test for the independence between nodes; for instance, selecting  $\alpha = 1$ , we would consider every node linked to each other. Thus we chose to maintain the typical threshold of 0.05.

So, while the Tabu Search algorithm seems to not be affected by hyperparameter tuning, the PC.Stable has the problem that its outcomes are particularly sensitive to changes of *alpha*. A possible solution to this issue can be found in [5]: here we define a function  $y_i = f(\theta) + \epsilon_i$ , with  $\epsilon_i$  being a Gaussian noise term, while  $f(\theta)$  can be seen as a black box objective function, where  $\theta$  depends on  $\alpha$ , the significance level, and  $T$ , the type of test. Each time we consider a new observation, a probabilistic model (typically a Gaussian process), is fitted to the data collected so far. In this way we obtain an optimization problem, where we want to maximize  $f$  in order to get the optimal value for  $\alpha$  and  $T$ .

A possible future improvement could be the implementation of such an approach, where the *alpha* value and the type of test are chosen together as the best parameters according to the data available.

## References

- [1] IBM HR Analytics Employee Attrition Performance. <https://www.kaggle.com/pavansubhasht/ibm-hr-analytics-attrition-dataset> 1.3
- [2] Beretta, S., Castelli, M., Gonçalves, I., Henriques, R., Ramazzotti, D.: Learning the structure of bayesian networks: A quantitative assessment of the effect of different algorithmic schemes. *Complexity* **2018** (2018) 1.1
- [3] Colombo, D., Maathuis, M.H.: A modification of the pc algorithm yielding order-independent skeletons. *arXiv preprint arXiv:1211.3295* (2012) 1.1
- [4] Colombo, D., Maathuis, M.H.: Order-independent constraint-based causal structure learning. *The Journal of Machine Learning Research* **15**(1), 3741–3782 (2014) 1.1
- [5] Córdoba, I., Garrido-Merchán, E.C., Hernández-Lobato, D., Bielza, C., Larranaga, P.: Bayesian optimization of the pc algorithm for learning gaussian bayesian networks. In: *Conference of the Spanish Association for Artificial Intelligence*. pp. 44–54. Springer (2018) 1.2, 4
- [6] Glover, F.: Tabu search—part i. *ORSA Journal on computing* **1**(3), 190–206 (1989) 1.1
- [7] Glover, F.: Tabu search—part ii. *ORSA Journal on computing* **2**(1), 4–32 (1990) 1.1